

Linear, Logistic, KNN, SVM, Random Forest

```
In [317]: import pandas as pd
import numpy as np
import seaborn as sns
import sklearn
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
```

```
In [318]: data = pd.read_csv("Insurance.csv")
data
```

Out[318]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [319]: data.columns
```

```
Out[319]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

```
In [320]: data['age'].unique()
```

```
Out[320]: array([19, 18, 28, 33, 32, 31, 46, 37, 60, 25, 62, 23, 56, 27, 52, 30, 34,
59, 63, 55, 22, 26, 35, 24, 41, 38, 36, 21, 48, 40, 58, 53, 43, 64,
20, 61, 44, 57, 29, 45, 54, 49, 47, 51, 42, 50, 39], dtype=int64)
```

```
In [321]: data['sex'].unique()
```

```
Out[321]: array(['female', 'male'], dtype=object)
```

```
In [322]: data['bmi'].unique()
```

```

Out[322]: array([27.9 , 33.77 , 33.   , 22.705, 28.88 , 25.74 , 33.44 , 27.74 ,
29.83 , 25.84 , 26.22 , 26.29 , 34.4 , 39.82 , 42.13 , 24.6 ,
30.78 , 23.845, 40.3 , 35.3 , 36.005, 32.4 , 34.1 , 31.92 ,
28.025, 27.72 , 23.085, 32.775, 17.385, 36.3 , 35.6 , 26.315,
28.6 , 28.31 , 36.4 , 20.425, 32.965, 20.8 , 36.67 , 39.9 ,
26.6 , 36.63 , 21.78 , 30.8 , 37.05 , 37.3 , 38.665, 34.77 ,
24.53 , 35.2 , 35.625, 33.63 , 28.   , 34.43 , 28.69 , 36.955,
31.825, 31.68 , 22.88 , 37.335, 27.36 , 33.66 , 24.7 , 25.935,
22.42 , 28.9 , 39.1 , 36.19 , 23.98 , 24.75 , 28.5 , 28.1 ,
32.01 , 27.4 , 34.01 , 29.59 , 35.53 , 39.805, 26.885, 38.285,
37.62 , 41.23 , 34.8 , 22.895, 31.16 , 27.2 , 26.98 , 39.49 ,
24.795, 31.3 , 38.28 , 19.95 , 19.3 , 31.6 , 25.46 , 30.115,
29.92 , 27.5 , 28.4 , 30.875, 27.94 , 35.09 , 29.7 , 35.72 ,
32.205, 28.595, 49.06 , 27.17 , 23.37 , 37.1 , 23.75 , 28.975,
31.35 , 33.915, 28.785, 28.3 , 37.4 , 17.765, 34.7 , 26.505,
22.04 , 35.9 , 25.555, 28.05 , 25.175, 31.9 , 36.   , 32.49 ,
25.3 , 29.735, 38.83 , 30.495, 37.73 , 37.43 , 24.13 , 37.145,
39.52 , 24.42 , 27.83 , 36.85 , 39.6 , 29.8 , 29.64 , 28.215,
37.   , 33.155, 18.905, 41.47 , 30.3 , 15.96 , 33.345, 37.7 ,
27.835, 29.2 , 26.41 , 30.69 , 41.895, 30.9 , 32.2 , 32.11 ,
31.57 , 26.2 , 30.59 , 32.8 , 18.05 , 39.33 , 32.23 , 24.035,
36.08 , 22.3 , 26.4 , 31.8 , 26.73 , 23.1 , 23.21 , 33.7 ,
33.25 , 24.64 , 33.88 , 38.06 , 41.91 , 31.635, 36.195, 17.8 ,
24.51 , 22.22 , 38.39 , 29.07 , 22.135, 26.8 , 30.02 , 35.86 ,
20.9 , 17.29 , 34.21 , 25.365, 40.15 , 24.415, 25.2 , 26.84 ,
24.32 , 42.35 , 19.8 , 32.395, 30.2 , 29.37 , 34.2 , 27.455,
27.55 , 20.615, 24.3 , 31.79 , 21.56 , 28.12 , 40.565, 27.645,
31.2 , 26.62 , 48.07 , 36.765, 33.4 , 45.54 , 28.82 , 22.99 ,
27.7 , 25.41 , 34.39 , 22.61 , 37.51 , 38.   , 33.33 , 34.865,
33.06 , 35.97 , 31.4 , 25.27 , 40.945, 34.105, 36.48 , 33.8 ,
36.7 , 36.385, 34.5 , 32.3 , 27.6 , 29.26 , 35.75 , 23.18 ,
25.6 , 35.245, 43.89 , 20.79 , 30.5 , 21.7 , 21.89 , 24.985,
32.015, 30.4 , 21.09 , 22.23 , 32.9 , 24.89 , 31.46 , 17.955,
30.685, 43.34 , 39.05 , 30.21 , 31.445, 19.855, 31.02 , 38.17 ,
20.6 , 47.52 , 20.4 , 38.38 , 24.31 , 23.6 , 21.12 , 30.03 ,
17.48 , 20.235, 17.195, 23.9 , 35.15 , 35.64 , 22.6 , 39.16 ,
27.265, 29.165, 16.815, 33.1 , 26.9 , 33.11 , 31.73 , 46.75 ,
29.45 , 32.68 , 33.5 , 43.01 , 36.52 , 26.695, 25.65 , 29.6 ,
38.6 , 23.4 , 46.53 , 30.14 , 30.   , 38.095, 28.38 , 28.7 ,
33.82 , 24.09 , 32.67 , 25.1 , 32.56 , 41.325, 39.5 , 34.3 ,
31.065, 21.47 , 25.08 , 43.4 , 25.7 , 27.93 , 39.2 , 26.03 ,
30.25 , 28.93 , 35.7 , 35.31 , 31.   , 44.22 , 26.07 , 25.8 ,
39.425, 40.48 , 38.9 , 47.41 , 35.435, 46.7 , 46.2 , 21.4 ,
23.8 , 44.77 , 32.12 , 29.1 , 37.29 , 43.12 , 36.86 , 34.295,
23.465, 45.43 , 23.65 , 20.7 , 28.27 , 35.91 , 29.   , 19.57 ,
31.13 , 21.85 , 40.26 , 33.725, 29.48 , 32.6 , 37.525, 23.655,
37.8 , 19.   , 21.3 , 33.535, 42.46 , 38.95 , 36.1 , 29.3 ,
39.7 , 38.19 , 42.4 , 34.96 , 42.68 , 31.54 , 29.81 , 21.375,
40.81 , 17.4 , 20.3 , 18.5 , 26.125, 41.69 , 24.1 , 36.2 ,
40.185, 39.27 , 34.87 , 44.745, 29.545, 23.54 , 40.47 , 40.66 ,
36.6 , 35.4 , 27.075, 28.405, 21.755, 40.28 , 30.1 , 32.1 ,
23.7 , 35.5 , 29.15 , 27.   , 37.905, 22.77 , 22.8 , 34.58 ,
27.1 , 19.475, 26.7 , 34.32 , 24.4 , 41.14 , 22.515, 41.8 ,
26.18 , 42.24 , 26.51 , 35.815, 41.42 , 36.575, 42.94 , 21.01 ,
24.225, 17.67 , 31.5 , 31.1 , 32.78 , 32.45 , 50.38 , 47.6 ,
25.4 , 29.9 , 43.7 , 24.86 , 28.8 , 29.5 , 29.04 , 38.94 ,
44.   , 20.045, 40.92 , 35.1 , 29.355, 32.585, 32.34 , 39.8 ,

```

```

24.605, 33.99 , 28.2 , 25. , 33.2 , 23.2 , 20.1 , 32.5 ,
37.18 , 46.09 , 39.93 , 35.8 , 31.255, 18.335, 42.9 , 26.79 ,
39.615, 25.9 , 25.745, 28.16 , 23.56 , 40.5 , 35.42 , 39.995,
34.675, 20.52 , 23.275, 36.29 , 32.7 , 19.19 , 20.13 , 23.32 ,
45.32 , 34.6 , 18.715, 21.565, 23. , 37.07 , 52.58 , 42.655,
21.66 , 32. , 18.3 , 47.74 , 22.1 , 19.095, 31.24 , 29.925,
20.35 , 25.85 , 42.75 , 18.6 , 23.87 , 45.9 , 21.5 , 30.305,
44.88 , 41.1 , 40.37 , 28.49 , 33.55 , 40.375, 27.28 , 17.86 ,
33.3 , 39.14 , 21.945, 24.97 , 23.94 , 34.485, 21.8 , 23.3 ,
36.96 , 21.28 , 29.4 , 27.3 , 37.9 , 37.715, 23.76 , 25.52 ,
27.61 , 27.06 , 39.4 , 34.9 , 22. , 30.36 , 27.8 , 53.13 ,
39.71 , 32.87 , 44.7 , 30.97 ])
```

In [323]: `data['children'].unique()`

Out[323]: `array([0, 1, 3, 2, 5, 4], dtype=int64)`

In [324]: `data['smoker'].unique()`

Out[324]: `array(['yes', 'no'], dtype=object)`

In [325]: `data['region'].unique()`

Out[325]: `array(['southwest', 'southeast', 'northwest', 'northeast'], dtype=object)`

In [326]: `data['charges'].unique()`

Out[326]: `array([16884.924 , 1725.5523, 4449.462 , ..., 1629.8335, 2007.945 ,
29141.3603])`

In [327]: `len(data['charges'])`

Out[327]: 1338

In [328]: `data.head()`

Out[328]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [329]: `data.tail()`

Out[329]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [330]: `len(data)`

Out[330]: 1338

In [331]: `data.index`

Out[331]: RangeIndex(start=0, stop=1338, step=1)

In [332]: `data.dtypes`

Out[332]:

age	int64
sex	object
bmi	float64
children	int64
smoker	object
region	object
charges	float64
dtype:	object

In [333]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [334]: `data.describe()`

Out[334]:

	age	bmi	children	charges
<b>count</b>	1338.000000	1338.000000	1338.000000	1338.000000
<b>mean</b>	39.207025	30.663397	1.094918	13270.422265
<b>std</b>	14.049960	6.098187	1.205493	12110.011237
<b>min</b>	18.000000	15.960000	0.000000	1121.873900
<b>25%</b>	27.000000	26.296250	0.000000	4740.287150
<b>50%</b>	39.000000	30.400000	1.000000	9382.033000
<b>75%</b>	51.000000	34.693750	2.000000	16639.912515
<b>max</b>	64.000000	53.130000	5.000000	63770.428010


In [335]: `corr = data.corr()`  
`corr`

Out[335]:

	age	bmi	children	charges
<b>age</b>	1.000000	0.109272	0.042469	0.299008
<b>bmi</b>	0.109272	1.000000	0.012759	0.198341
<b>children</b>	0.042469	0.012759	1.000000	0.067998
<b>charges</b>	0.299008	0.198341	0.067998	1.000000

In [336]: `sns.heatmap(corr)`

Out[336]: `<AxesSubplot:>`



In [337]: `data.isna()`

Out[337]:

	age	sex	bmi	children	smoker	region	charges
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...
1333	False	False	False	False	False	False	False
1334	False	False	False	False	False	False	False
1335	False	False	False	False	False	False	False
1336	False	False	False	False	False	False	False
1337	False	False	False	False	False	False	False

In [338]: `data.isna().sum()`

```
Out[338]: age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
```

In [339]: `data.isnull()`

Out[339]:

	age	sex	bmi	children	smoker	region	charges
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...
1333	False	False	False	False	False	False	False
1334	False	False	False	False	False	False	False
1335	False	False	False	False	False	False	False
1336	False	False	False	False	False	False	False
1337	False	False	False	False	False	False	False

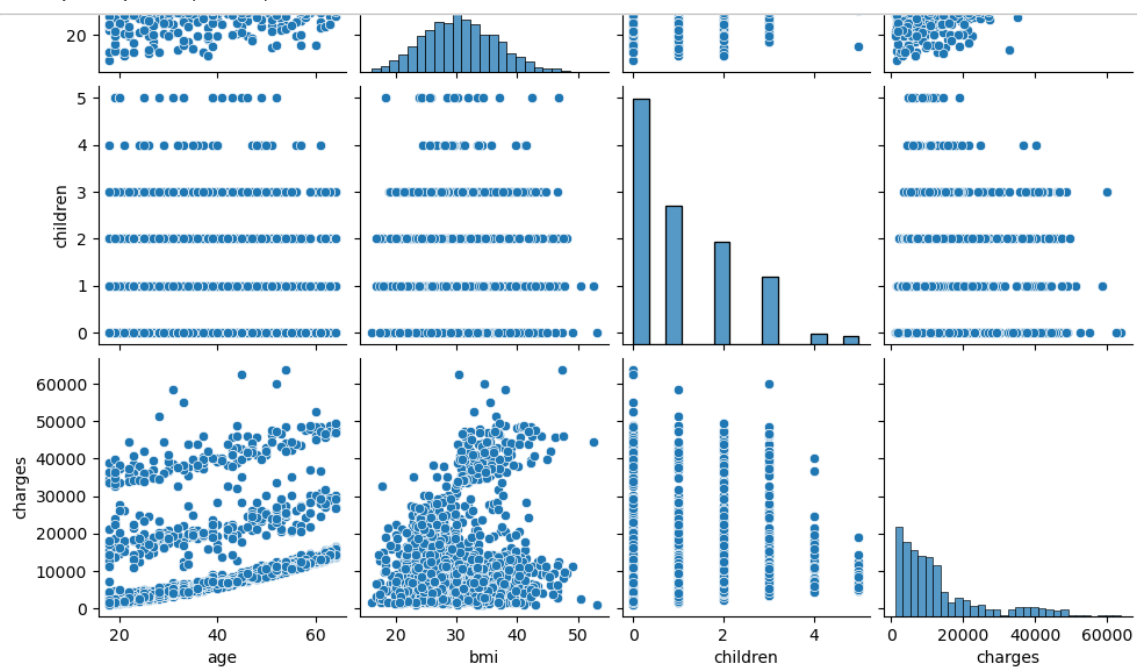
1338 rows × 7 columns

In [340]: `data.isnull().sum()`

```
Out[340]: age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
dtype: int64
```




```
In [341]: sns.pairplot(data)
```




```
In [342]: sns.histplot(x=data['age'],data=data)
```

```
Out[342]: <AxesSubplot:xlabel='age', ylabel='Count'>
```




```
In [343]: sns.histplot(x='bmi', data=data)
```

```
Out[343]: <AxesSubplot:xlabel='bmi', ylabel='Count'>
```




```
In [344]: sns.histplot(x=data['children'],data=data)
```

```
Out[344]: <AxesSubplot:xlabel='children', ylabel='Count'>
```




```
In [345]: sns.histplot(x=data['charges'],data=data)
```

```
Out[345]: <AxesSubplot:xlabel='charges', ylabel='Count'>
```




```
In [346]: sns.countplot(x = 'sex', data = data, hue = 'children')
```

```
Out[346]: <AxesSubplot:xlabel='sex', ylabel='count'>
```




```
In [347]: sns.countplot(x='region',data=data,hue='sex')
```


```
Out[347]: <AxesSubplot:xlabel='region', ylabel='count'>
```




```
In [348]: plt.scatter(x=data['bmi'],y=data['charges'],c='r')  
plt.xlabel('bmi')  
plt.ylabel('charges')  
plt.show()
```




```
In [349]: plt.bar(data['bmi'],data['charges'])  
plt.xlabel('bmi')  
plt.ylabel('charges')  
plt.show()
```



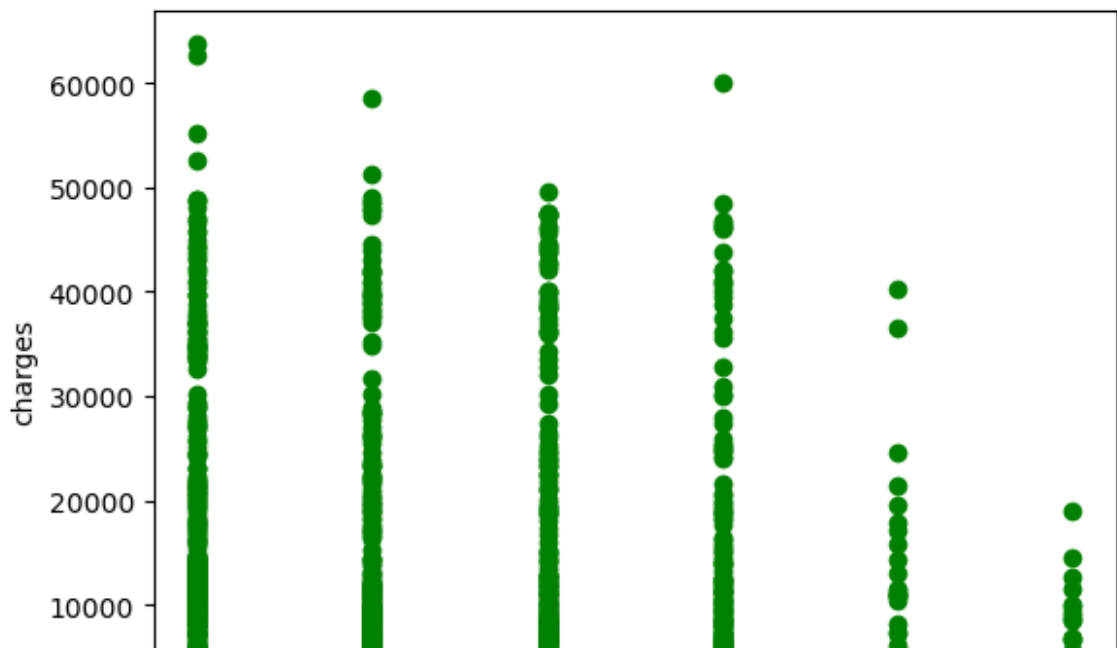
```
In [350]: plt.scatter(x=data['age'],y=data['charges'],c='y')  
plt.xlabel('age')  
plt.ylabel('charges')  
plt.show()
```



```
In [351]: plt.bar(data['age'],data['charges'])  
plt.xlabel('age')  
plt.ylabel('charges')  
plt.show()
```




```
In [352]: plt.scatter(x=data['children'],y=data['charges'],c='g')  
plt.xlabel('children')  
plt.ylabel('charges')  
plt.show()
```




```
In [353]: sns.boxplot(x=data["children"], y=data["charges"], data=data)
```

```
Out[353]: <AxesSubplot:xlabel='children', ylabel='charges'>
```






```
In [354]: plt.scatter(x=data['children'],y=data['bmi'],c='g')
plt.xlabel('children')
plt.ylabel('bmi')
plt.show()
```




```
In [355]: sns.boxplot(x=data["children"], y=data["bmi"], data=data)
```

```
Out[355]: <AxesSubplot:xlabel='children', ylabel='bmi'>
```




```
In [356]: sns.boxplot(x=data["region"], y=data["bmi"], data=data)
```

```
Out[356]: <AxesSubplot:xlabel='region', ylabel='bmi'>
```



```
In [357]: sns.boxplot(x=data["region"], y=data["charges"], data=data)
```

```
Out[357]: <AxesSubplot:xlabel='region', ylabel='charges'>
```



```
In [358]: print(f"Our data has dimensions: {data.shape[0]} by {data.shape[1]}")
```

Our data has dimensions: 1338 by 7

```
In [359]: x = data['age']  
y = data['charges']
```

```
In [360]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_stat
```

```
In [361]: print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)
```

(1070,)  
(268,)  
(1070,)  
(268,)

```
In [362]: x_train=x_train.values.reshape(-1,1)
```

```
In [363]: x_test=x_test.values.reshape(-1,1)
```

```
In [364]: model = LinearRegression()
```

```
In [365]: model.fit(x_train,y_train)
```

```
Out[365]: LinearRegression()
```

```
In [366]: model.score(x_test,y_test)
```

```
Out[366]: 0.07899146595873285
```

```
In [367]: print(f"intercept: {model.intercept_}")
```

```
intercept: 2998.5777306780055
```

```
In [368]: print(f"slope: {model.coef_}")
```

```
slope: [262.31687345]
```

```
In [369]: y_pred = model.intercept_ + model.coef_*x_test  
print(f"y_pred : {y_pred}")
```

```
y_pred : [[ 7982.59832619]
 [17950.63951722]
 [16376.73827653]
 [15852.10452963]
 [ 8507.23207309]
 [11130.40080757]
 [15852.10452963]
 [17688.32264377]
 [10081.13331377]
 [12704.30204826]
 [18737.59013756]
 [17163.68889687]
 [13228.93579515]
 [14015.8864155 ]
 [ 7720.28145274]
 [15589.78765618]
 [11655.03455446]
 [12179.66830136]
 [18475.27326411]
 [19262.22388446]
 [16639.05514998]
 [10868.08393412]
 [15065.15390929]
 [16114.42140308]
 [ 7720.28145274]
 [13753.56954205]
 [14278.20328894]
 [15589.78765618]
 [11655.03455446]
 [10868.08393412]
 [17688.32264377]
 [11917.35142791]
 [16901.37202342]
 [ 7982.59832619]
 [15589.78765618]
 [16901.37202342]
 [10868.08393412]
 [ 8507.23207309]
 [18212.95639066]
 [14802.83703584]
 [10343.45018722]
 [18212.95639066]
 [16114.42140308]
 [18475.27326411]
 [11130.40080757]
 [18737.59013756]
 [ 8244.91519964]
 [10343.45018722]
 [16376.73827653]
 [17950.63951722]
 [18999.90701101]
 [17426.00577032]
 [ 9294.18269343]
 [14802.83703584]
 [ 7720.28145274]
 [11392.71768101]
 [14278.20328894]
```

```
[18737.59013756]
[ 9031.86581998]
[ 7720.28145274]
[10605.76706067]
[18999.90701101]
[ 9031.86581998]
[ 9031.86581998]
[18737.59013756]
[15589.78765618]
[15852.10452963]
[17950.63951722]
[ 7982.59832619]
[ 9294.18269343]
[16639.05514998]
[11130.40080757]
[ 7982.59832619]
[17163.68889687]
[19262.22388446]
[14278.20328894]
[ 9031.86581998]
[10605.76706067]
[17688.32264377]
[16114.42140308]
[ 8507.23207309]
[17426.00577032]
[16901.37202342]
[ 9294.18269343]
[13491.2526686 ]
[ 9556.49956688]
[14015.8864155 ]
[14540.52016239]
[15852.10452963]
[13753.56954205]
[12179.66830136]
[15852.10452963]
[18737.59013756]
[15589.78765618]
[ 7982.59832619]
[16901.37202342]
[16376.73827653]
[17688.32264377]
[18212.95639066]
[12179.66830136]
[17950.63951722]
[ 9556.49956688]
[10868.08393412]
[ 7982.59832619]
[19786.85763135]
[ 9031.86581998]
[ 7982.59832619]
[14802.83703584]
[ 7720.28145274]
[17426.00577032]
[11917.35142791]
[ 9818.81644033]
[ 7982.59832619]
[19524.5407579 ]
```



```
[ 9294.18269343]
[ 7982.59832619]
[11655.03455446]
[19524.5407579 ]
[12179.66830136]
[ 8244.91519964]
[16114.42140308]
[15589.78765618]
[12966.6189217 ]
[15065.15390929]
[17688.32264377]
[15589.78765618]
[17426.00577032]
[12179.66830136]
[14015.8864155 ]
[14540.52016239]
[ 7720.28145274]
[ 7982.59832619]
[12966.6189217 ]
[11655.03455446]
[15065.15390929]
[ 7982.59832619]
[ 7982.59832619]
[ 7720.28145274]
[16639.05514998]
[14802.83703584]
[19262.22388446]
[12179.66830136]
[11917.35142791]
[19786.85763135]
[18212.95639066]
[17426.00577032]
[ 7720.28145274]
[13753.56954205]
[ 7720.28145274]
[12704.30204826]
[14540.52016239]
[ 9818.81644033]
[16639.05514998]
[17426.00577032]
[15589.78765618]
[14015.8864155 ]
[16376.73827653]
[16376.73827653]
[15589.78765618]
[ 9031.86581998]
[11655.03455446]
[19786.85763135]
[18212.95639066]
[ 9818.81644033]
[11392.71768101]
[ 8507.23207309]
[12441.98517481]
[17163.68889687]
[17688.32264377]
[12441.98517481]
[11392.71768101]
```

```
[10343.45018722]
[17688.32264377]
[18475.27326411]
[10081.13331377]
[12966.6189217 ]
[ 9031.86581998]
[18475.27326411]
[10343.45018722]
[12179.66830136]
[18737.59013756]
[18999.90701101]
[15852.10452963]
[14015.8864155 ]
[ 9556.49956688]
[10081.13331377]
[ 8769.54894653]
[15327.47078274]
[15065.15390929]
[13753.56954205]
[18999.90701101]
[16639.05514998]
[11130.40080757]
[ 7982.59832619]
[12966.6189217 ]
[16376.73827653]
[12441.98517481]
[ 7982.59832619]
[16376.73827653]
[15065.15390929]
[13753.56954205]
[13753.56954205]
[17163.68889687]
[13228.93579515]
[12966.6189217 ]
[ 7720.28145274]
[13491.2526686 ]
[14015.8864155 ]
[ 8244.91519964]
[17950.63951722]
[10868.08393412]
[16901.37202342]
[ 9294.18269343]
[ 9818.81644033]
[11392.71768101]
[19786.85763135]
[17426.00577032]
[ 8769.54894653]
[ 7720.28145274]
[14540.52016239]
[ 8244.91519964]
[16901.37202342]
[17426.00577032]
[10081.13331377]
[ 7982.59832619]
[ 9294.18269343]
[11392.71768101]
[ 9031.86581998]
```

```
[ 7982.59832619]
[19524.5407579 ]
[ 8507.23207309]
[18475.27326411]
[16376.73827653]
[11917.35142791]
[17163.68889687]
[17426.00577032]
[10868.08393412]
[ 7982.59832619]
[15065.15390929]
[11392.71768101]
[15327.47078274]
[ 8507.23207309]
[19262.22388446]
[ 9294.18269343]
[17426.00577032]
[16639.05514998]
[14278.20328894]
[15327.47078274]
[ 9556.49956688]
[19786.85763135]
[16901.37202342]
[11130.40080757]
[18737.59013756]
[16376.73827653]
[ 9818.81644033]
[ 8769.54894653]
[ 7720.28145274]
[13491.2526686 ]
[ 7982.59832619]
[18999.90701101]
[14540.52016239]
[14015.8864155 ]
[15065.15390929]
[13228.93579515]
[14540.52016239]
[15327.47078274]
[ 7982.59832619]
[12179.66830136]]
```

```
In [370]: print(x_train.shape)
          print(y_test.shape)
          print(y_pred.shape)
```

```
(1070, 1)
(268,)
(268, 1)
```

```
In [371]: y_test = y_test.values.reshape(-1,1)
```

```
In [372]: print(y_test.shape)
```

```
(268, 1)
```

```
In [373]: rss = np.sum((y_test - y_pred) ** 2)
```

```
In [374]: print("RSS:", rss)
```

```
RSS: 36846983136.88877
```

```
In [375]: x1 = data['bmi']  
y = data['charges']
```

```
In [376]: x1_train, x1_test, y_train, y_test = train_test_split(x1, y, test_size = 0.2,
```

```
In [377]: print(x1_train.shape)  
print(x1_test.shape)  
print(y_train.shape)  
print(y_test.shape)
```

```
(1070,)  
(268,)  
(1070,)  
(268,)
```

```
In [378]: x1_train = x1_train.values.reshape(-1,1)
```

```
In [379]: x1_test = x1_test.values.reshape(-1,1)
```

```
In [380]: model1 = LinearRegression()
```

```
In [381]: model1.fit(x1_train,y_train)
```

```
Out[381]: LinearRegression()
```

```
In [382]: model1.score(x1_test,y_test)
```

```
Out[382]: 0.05716373981200229
```

```
In [383]: print(f"intercept: {model1.intercept_}")  
print(f"slope: {model1.coef_}")
```

```
intercept: 2001.129553376486  
slope: [262.31687345]
```

```
In [384]: y1_pred = model1.coef_*x1_test + model1.intercept_  
print(f"y1_pred: {y1_pred}")
```

```
y1_pred: [[15009.47211896]
[13548.64220518]
[15547.67261351]
[15412.20718291]
[10165.667668 ]
[14592.09214359]
[15734.39523407]
[12470.41060215]
[11531.30565757]
[11666.77108818]
[10982.12147946]
[19099.06363198]
[13841.54043351]
[12618.69033024]
[14606.73705501]
[13058.03767273]
[10313.94739609]
[12142.73070921]
[15371.93367651]
[14156.40602896]
[14174.71216823]
[16298.2243236 ]
[12157.37562062]
[11739.99564526]
[13026.91723597]
[12400.84727292]
[12018.24896217]
[14204.00199106]
[14046.56919333]
[11322.61566989]
[15108.32527102]
[11886.44475942]
[10279.16573148]
[13687.76886363]
[11465.4035562 ]
[10905.23569452]
[14925.26387832]
[11410.48513839]
[16298.2243236 ]
[13166.04389443]
[14083.18147188]
[14592.09214359]
[12190.32667131]
[15584.28489205]
[10641.62728903]
[12470.41060215]
[14928.92510617]
[13599.89939514]
[ 8609.64583001]
[12539.97393138]
[13559.62588874]
[12874.97628003]
[10661.76404223]
[11410.48513839]
[12018.24896217]
[18191.07912416]
[13922.0874463 ]]
```

```
[14852.03932123]
[10870.45402991]
[17748.07055382]
[15009.47211896]
[14279.05716207]
[12431.96770968]
[12296.50227908]
[15531.19708817]
[13826.89552209]
[12505.19226676]
[10780.75394749]
[11009.58068837]
[ 9513.96910997]
[14626.87380821]
[15948.57706353]
[12582.0780517 ]
[11009.58068837]
[12992.13557136]
[14800.78213128]
[12018.24896217]
[12230.60017771]
[11264.03602423]
[11600.8689868 ]
[ 8157.48419003]
[14279.05716207]
[11079.14401759]
[12728.52716586]
[14000.80384516]
[18674.3612009 ]
[11264.03602423]
[18070.25860498]
[14209.49383284]
[14105.148839 ]
[12261.72061447]
[10731.32737146]
[14123.45497827]
[13801.26692711]
[10279.16573148]
[12794.42926724]
[11447.09741693]
[15734.39523407]
[15357.2887651 ]
[10592.200713 ]
[14452.96548514]
[10835.6723653 ]
[10934.52551736]
[14742.20248561]
[10418.29238993]
[12366.06560831]
[11044.36235298]
[13058.03767273]
[15653.84822128]
[12922.57224213]
[12150.05316492]
[12955.52329282]
[11218.27067605]
[17131.15366041]
```

```
[11809.55897448]  
[12508.85349462]  
[17524.73565472]  
[11183.48901144]  
[16539.86536197]  
[13687.76886363]  
[11373.87285985]  
[15044.25378357]  
[13237.43783758]  
[10165.667668 ]  
[13640.17290153]  
[14559.14109291]  
[11287.83400528]  
[10941.84797307]  
[16781.50640034]  
[15584.28489205]  
[15976.03627244]  
[ 9653.09576843]  
[12192.15728524]  
[15090.01913175]  
[12106.11843067]  
[ 9543.2589328 ]  
[11357.3973345 ]  
[16016.30977883]  
[ 8713.99082385]  
[16018.14039276]  
[13618.20553441]  
[14566.46354861]  
[11264.03602423]  
[15170.56614454]  
[13652.98719902]  
[12122.59395601]  
[12754.15576084]  
[13936.73235771]  
[14888.65159977]  
[10557.41904839]  
[15935.76276604]  
[13305.17055288]  
[15739.88707585]  
[13861.6771867 ]  
[16700.95938755]  
[11849.83248088]  
[13826.89552209]  
[12157.37562062]  
[12018.24896217]  
[14595.75337145]  
[13076.343812 ]  
[15533.02770209]  
[13720.71991432]  
[12984.81311565]  
[14139.93050361]  
[15492.7541957 ]  
[14888.65159977]  
[12691.91488732]  
[11959.6693165 ]  
[12157.37562062]  
[12911.58855857]
```



```
[14731.21880205]
[ 9305.27912229]
[10661.76404223]
[ 9861.78575611]
[11983.46729755]
[14888.65159977]
[13861.6771867 ]
[ 8331.3925131 ]
[17868.891073 ]
[12853.0089129 ]
[10070.47574379]
[12957.35390674]
[17131.15366041]
[10244.38406686]
[12673.60874805]
[12270.8736841 ]
[18593.81418811]
[ 9270.49745767]
[12853.0089129 ]
[13826.89552209]
[17305.06198347]
[ 9470.03437572]
[10974.79902375]
[ 9131.37079922]
[ 9543.2589328 ]
[12633.33524166]
[14870.3454605 ]
[10940.01735914]
[12270.8736841 ]
[14526.19004222]
[15584.28489205]
[10421.95361778]
[14485.91653582]
[12713.88225445]
[11787.59160736]
[14487.74714975]
[11586.22407539]
[16818.11867888]
[12508.85349462]
[12226.93894985]
[10378.01888354]
[14313.83882668]
[14522.52881437]
[14888.65159977]
[13460.77273668]
[12609.5372606 ]
[13358.25835677]
[14324.82251024]
[13058.03767273]
[13200.82555904]
[ 9653.09576843]
[14000.80384516]
[ 8922.68081153]
[12192.15728524]
[10568.40273195]
[12887.79057751]
[14229.63058604]
```

```
[12362.40438045]
[15492.7541957 ]
[14313.83882668]
[12190.32667131]
[17707.79704742]
[ 8957.46247615]
[12574.75559599]
[12948.20083711]
[10418.29238993]
[13200.82555904]
[16435.52036813]
[12995.79679921]
[19399.28431601]
[16459.31834918]
[15496.41542355]
[13720.71991432]
[14083.18147188]
[10905.23569452]
[14928.92510617]
[13717.05868647]
[13801.26692711]
[15879.01373431]
[15705.10541124]
[11844.3406391 ]
[11425.13004981]
[14204.00199106]
[12783.44558367]
[16574.64702658]
[14279.05716207]
[12296.50227908]
[ 9479.18744536]
[15322.50710048]
[12915.24978642]
[11635.65065142]
[14244.27549745]
[14925.26387832]
[16261.61204506]
[16257.9508172 ]
[15392.07042971]
[11566.08732219]]
```


```
In [385]: y_test = y_test.values.reshape(-1,1)
```

```
In [386]: RSS1 = np.sum((y_test - y1_pred) ** 2)
```

```
In [387]: print(f"RSS :{RSS1}")
```

```
RSS :37720249591.561134
```

```
In [388]: plt.scatter(x=data['charges'],y=data['smoker'])
plt.xlabel("charges")
plt.ylabel("smoker")
plt.title("charges vs smoker")
plt.show()
```



```
In [389]: smoke = pd.get_dummies(data['smoker'],drop_first=True)
```

```
In [390]: data['smoke'] = smoke
```

```
In [391]: data.head()
```

Out[391]:

	age	sex	bmi	children	smoker	region	charges	smoke
0	19	female	27.900	0	yes	southwest	16884.92400	1
1	18	male	33.770	1	no	southeast	1725.55230	0
2	28	male	33.000	3	no	southeast	4449.46200	0
3	33	male	22.705	0	no	northwest	21984.47061	0
4	32	male	28.880	0	no	northwest	3866.85520	0

```
In [392]: xa = data['charges']  
ya = data['smoke']
```

```
In [393]: xa_train, xa_test, ya_train, ya_test = train_test_split(xa,ya,test_size=0.2,ra
```

```
In [394]: print(xa_train.shape)  
print(xa_test.shape)  
print(ya_train.shape)  
print(ya_test.shape)
```

```
(1070,)  
(268,)  
(1070,)  
(268,)
```

```
In [395]: xa_train = xa_train.values.reshape(-1,1)  
xa_test = xa_test.values.reshape(-1,1)
```

```
In [396]: modela = LogisticRegression()  
modela.fit(xa_train,ya_train)
```

```
Out[396]: LogisticRegression()
```

```
In [397]: modela.score(xa_test,ya_test)
```

```
Out[397]: 0.9104477611940298
```

```
In [398]: print(f"intercept: {modela.intercept_}")  
print(f"slope: {modela.coef_}")
```

```
intercept: [-5.65519559]  
slope: [[0.00025126]]
```

```
In [399]: modela.classes_
```

```
Out[399]: array([0, 1], dtype=uint8)
```

```
In [400]: modela.predict_proba(xa_test)
```

```
Out[400]: array([[9.94735580e-01, 5.26442035e-03],  
 [9.42812396e-01, 5.71876039e-02],  
 [9.69064287e-01, 3.09357129e-02],  
 [9.54639685e-01, 4.53603154e-02],  
 [9.94099298e-01, 5.90070182e-03],  
 [1.66198820e-02, 9.83380118e-01],  
 [9.65015950e-01, 3.49840505e-02],  
 [9.38539065e-01, 6.14609354e-02],  
 [9.92487229e-01, 7.51277130e-03],  
 [6.78235221e-01, 3.21764779e-01],  
 [9.22852359e-01, 7.71476405e-02],  
 [9.40251179e-01, 5.97488210e-02],  
 [9.83087762e-01, 1.69122385e-02],  
 [9.79838602e-01, 2.01613981e-02],  
 [9.95364614e-01, 4.63538622e-03],  
 [9.67759830e-01, 3.22401699e-02],  
 [5.32771446e-01, 4.67228554e-01],  
 [9.82766883e-01, 1.72331167e-02],  
 [1.89608694e-01, 8.10391306e-01],  
 [9.96577886e-01, 9.24228128e-03]
```

```
In [401]: ya_pred = modela.predict(x_test)
```

```
In [402]: print(f"ya_pred :{ya_pred}")
```

[illegible]

```
In [403]: pd.DataFrame(confusion_matrix(ya_test,ya_pred),columns=['Predicted No', 'Predicted Yes'])
```

Out[403]:

	Predicted No	Predicted Yes
No	214	0
Yes	54	0

```
In [404]: print(classification_report(ya_test,ya_pred))
```

	precision	recall	f1-score	support
0	0.80	1.00	0.89	214
1	0.00	0.00	0.00	54
accuracy			0.80	268
macro avg	0.40	0.50	0.44	268
weighted avg	0.64	0.80	0.71	268

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

```
In [405]: gender = pd.get_dummies(data['sex'],drop_first=True)
```

```
In [406]: data['gender'] = gender
```

```
In [407]: data.head()
```

Out[407]:

	age	sex	bmi	children	smoker	region	charges	smoke	gender
0	19	female	27.900	0	yes	southwest	16884.92400	1	0
1	18	male	33.770	1	no	southeast	1725.55230	0	1
2	28	male	33.000	3	no	southeast	4449.46200	0	1
3	33	male	22.705	0	no	northwest	21984.47061	0	1
4	32	male	28.880	0	no	northwest	3866.85520	0	1

```
In [408]: xb = data['charges']
          yb = data['gender']
```

```
In [409]: xb_train, xb_test, yb_train, yb_test = train_test_split(xb, yb, test_size = 0.
```



```
In [418]: yb_pred = modelb.predict(xb_test)
```

```
In [419]: print(f"yb_pred:{yb_pred}")
```

[illegible]

```
In [420]: pd.DataFrame(confusion_matrix(yb_test,yb_pred),columns=['Predicted No', 'Predicted Yes'])
```

Out[420]:

	Predicted No	Predicted Yes
No	0	137
Yes	0	131

```
In [421]: print(classification_report(yb_test, yb_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	137
1	0.49	1.00	0.66	131
accuracy			0.49	268
macro avg	0.24	0.50	0.33	268
weighted avg	0.24	0.49	0.32	268

```
C:\Users\Noah\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division` parameter
to control this behavior.
```

```
warn_prf(average, modifier, msg_start, len(result))
```

```
C:\Users\Noah\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division` parameter
to control this behavior.
```

```
warn prf(average, modifier, msg start, len(result))
```

```
C:\Users\Noah\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:
1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being
set to 0.0 in labels with no predicted samples. Use `zero_division` parameter
to control this behavior.
```

```
warn prf(average, modifier, msg start, len(result))
```

```
In [422]: xc = data['bmi']  
          yc = data['gender']
```



```
In [423]: xc_train, xc_test, yc_train, yc_test = train_test_split(xc, yc, test_size=0.2,
```

```
In [424]: print(xc_train.shape)
print(xc_test.shape)
print(yc_train.shape)
print(yc_test.shape)
```

```
(1070,)
(268,)
(1070,)
(268,)
```

```
In [425]: xc_train = xc_train.values.reshape(-1,1)
xc_test = xc_test.values.reshape(-1,1)
```

```
In [426]: modelc = LogisticRegression()
```

```
In [427]: modelc.fit(xc_train,yc_train)
```

```
Out[427]: LogisticRegression()
```

```
In [428]: modelc.score(xc_test,yc_test)
```

```
Out[428]: 0.5223880597014925
```

```
In [429]: modelc.classes_
```

```
Out[429]: array([0, 1], dtype=uint8)
```

```
In [430]: print(f"modelc intercept: {modelc.intercept_}")
```

```
modelc intercept: [-0.33149167]
```

```
In [431]: print(f"modelc slope: {modelc.coef_}")
```

```
modelc slope: [[0.01202926]]
```



```
In [437]: from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
data['region'] = label_encoder.fit_transform(data['region'])
print(data.head())
```

	age	sex	bmi	children	smoker	region	charges	smoke	gender
0	19	female	27.900	0	yes	3	16884.92400	1	0
1	18	male	33.770	1	no	2	1725.55230	0	1
2	28	male	33.000	3	no	2	4449.46200	0	1
3	33	male	22.705	0	no	1	21984.47061	0	1
4	32	male	28.880	0	no	1	3866.85520	0	1

```
In [438]: xk = data['bmi']
yk = data['region']
xk
yk
```

```
Out[438]: 0      3
1      2
2      2
3      1
4      1
..
1333    1
1334    0
1335    2
1336    3
1337    1
Name: region, Length: 1338, dtype: int32
```

```
In [439]: xk_train,xk_test,yk_train,yk_test = train_test_split(xk,yk,test_size=0.2,randc
```

```
In [440]: print(xk_train.shape)
print(xk_test.shape)
print(yk_train.shape)
print(yk_test.shape)
```

```
(1070,)
(268,)
(1070,)
(268,)
```

```
In [441]: xk_train = xk_train.values.reshape(-1,1)
xk_test = xk_test.values.reshape(-1,1)
```

```
In [442]: NN = KNeighborsClassifier(n_neighbors=5)
```

```
In [443]: NN.fit(xk_train,yk_train)
```

```
Out[443]: KNeighborsClassifier()
```

In [444]: `NN.score(xk_test,yk_test)`

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

Out[444]: 0.40298507462686567

In [445]: `yNN_pred = NN.predict(xk_test)`

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

In [446]: `print(f"yNN_pred : {yNN_pred}")`

```
yNN_pred : [2 3 1 1 0 3 3 1 0 1 3 2 0 0 2 3 0 2 1 0 0 2 0 1 1 3 0 2 0 0 1 0 1
 0 1 1 0
 2 2 1 2 3 1 3 0 1 3 0 0 0 3 2 1 2 0 3 0 2 3 2 2 2 0 1 1 1 3 0 3 3 2 2 3 3
 1 0 0 1 0 1 0 2 0 2 1 2 0 3 0 3 0 3 2 3 1 2 3 3 1 0 2 1 1 3 0 1 1 3 2 0 2
 2 0 0 2 3 2 0 3 0 3 1 0 0 2 2 1 1 0 3 0 2 1 2 3 3 3 0 3 0 0 2 0 2 0 1 3 2
 2 0 2 1 2 0 2 1 1 0 0 3 1 1 0 1 0 2 2 2 1 0 3 3 1 1 3 0 2 0 0 3 0 0 2 0 0
 1 0 2 0 0 1 2 0 2 3 3 1 2 1 0 2 3 0 3 2 3 1 1 1 3 1 0 1 3 2 1 0 2 2 3 1 2
 1 0 1 0 1 3 1 2 1 1 2 0 1 2 0 1 3 1 2 2 2 0 2 1 3 0 3 3 3 1 2 2 2 2 1 0
 3 2 1 3 0 2 2 1 0]
```

In [447]: `pd.DataFrame(confusion_matrix(yk_test,yNN_pred),columns=['Predicted 0','Predicted 1','Predicted 2','Predicted 3'])`

Out[447]:

	Predicted 0	Predicted 1	Predicted 2	Predicted 3
Actual 0	29	20	6	6
Actual 1	17	23	11	15
Actual 2	16	14	37	15
Actual 3	15	10	15	19

```
In [448]: print(classification_report(yk_test,yNN_pred))
```

	precision	recall	f1-score	support
0	0.38	0.48	0.42	61
1	0.34	0.35	0.35	66
2	0.54	0.45	0.49	82
3	0.35	0.32	0.33	59
accuracy			0.40	268
macro avg	0.40	0.40	0.40	268
weighted avg	0.41	0.40	0.40	268

```
In [449]: xk1 = data['charges']
xk1
```

```
Out[449]: 0      16884.92400
1      1725.55230
2      4449.46200
3      21984.47061
4      3866.85520
...
1333    10600.54830
1334     2205.98080
1335     1629.83350
1336     2007.94500
1337     29141.36030
Name: charges, Length: 1338, dtype: float64
```

```
In [450]: yk = data['region']
yk
```

```
Out[450]: 0      3
1      2
2      2
3      1
4      1
..
1333    1
1334    0
1335    2
1336    3
1337    1
Name: region, Length: 1338, dtype: int32
```

```
In [451]: xk1_train, xk1_test, yk_train, yk_test = train_test_split(xk1,yk,test_size=0.2)
```

```
In [452]: print(xk1_test.shape)
print(xk1_train.shape)
print(yk_test.shape)
print(yk_train.shape)
```

```
(268,)
(1070,)
(268,)
(1070,)
```

```
In [453]: xk1_test = xk1_test.values.reshape(-1,1)
xk1_train = xk1_train.values.reshape(-1,1)
```

```
In [454]: NN1 = KNeighborsClassifier(n_neighbors=5)
```

```
In [455]: NN1.fit(xk1_train,yk_train)
```

```
Out[455]: KNeighborsClassifier()
```

```
In [456]: NN1.score(xk1_test,yk_test)
```

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
Out[456]: 0.332089552238806
```

```
In [457]: yNN1_pred = NN1.predict(xk1_test)
```

C:\Users\Noah\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
In [458]: print(f"yNN1_pred : {yNN1_pred}")
```

```
yNN1_pred : [2 3 3 0 1 0 0 0 0 2 1 0 1 0 2 1 1 1 1 2 0 0 1 1 0 0 0 3 1 3 0 0
0 0 1 0 2
0 3 2 0 0 0 0 1 1 0 3 3 0 1 2 0 2 2 1 0 2 3 2 2 0 1 1 0 3 1 2 1 0 2 2 3 1
1 0 1 1 1 0 1 1 3 3 0 2 0 2 0 1 0 1 0 0 2 0 3 1 1 1 3 3 1 3 2 0 3 1 0 1 1
2 2 0 1 3 0 1 2 0 0 3 2 3 1 0 3 2 1 1 3 2 3 2 2 3 2 2 0 0 0 3 3 1 0 1 2 0
2 0 2 0 0 2 2 2 0 3 0 2 2 2 0 2 1 2 0 0 1 2 1 3 3 0 2 1 0 1 2 2 0 0 1 2 3
1 3 2 1 1 2 2 0 1 0 0 1 1 0 2 3 3 1 0 2 2 3 0 3 1 1 0 2 2 0 2 3 3 2 0 2 3
1 2 2 3 0 3 2 1 3 0 3 0 1 1 1 2 0 1 3 1 1 3 0 0 1 0 1 3 2 0 1 2 2 0 3 3 2
2 2 0 2 1 2 2 1 3]
```

```
In [459]: pd.DataFrame(confusion_matrix(yk_test,yNN1_pred),columns=['Predicted 0','Predi
```

Out[459]:

	Predicted 0	Predicted 1	Predicted 2	Predicted 3
Actual 0	25	11	12	13
Actual 1	20	22	14	10
Actual 2	18	21	30	13
Actual 3	18	16	13	12

```
In [460]: print(classification_report(yk_test,yNN1_pred))
```

	precision	recall	f1-score	support
0	0.31	0.41	0.35	61
1	0.31	0.33	0.32	66
2	0.43	0.37	0.40	82
3	0.25	0.20	0.22	59
accuracy			0.33	268
macro avg	0.33	0.33	0.32	268
weighted avg	0.34	0.33	0.33	268

KNN not good on large datasets

```
In [461]: data.head()
```

Out[461]:

	age	sex	bmi	children	smoker	region	charges	smoke	gender
0	19	female	27.900	0	yes	3	16884.92400	1	0
1	18	male	33.770	1	no	2	1725.55230	0	1
2	28	male	33.000	3	no	2	4449.46200	0	1
3	33	male	22.705	0	no	1	21984.47061	0	1
4	32	male	28.880	0	no	1	3866.85520	0	1

```
In [462]: from sklearn import svm
```

```
In [463]: xs = data['charges']
```

```
In [464]: ys = data['smoke']
```

In [465]: `print(xs)`

```
0      16884.92400
1      1725.55230
2      4449.46200
3      21984.47061
4      3866.85520
...
1333   10600.54830
1334    2205.98080
1335    1629.83350
1336    2007.94500
1337   29141.36030
Name: charges, Length: 1338, dtype: float64
```

In [466]: `print(ys)`

```
0      1
1      0
2      0
3      0
4      0
..
1333   0
1334   0
1335   0
1336   0
1337   1
Name: smoke, Length: 1338, dtype: uint8
```

In [467]: `xs_train, xs_test, ys_train, ys_test = train_test_split(xs,ys,test_size=0.2,ra`

In [468]: `print(xs_train.shape)`  
`print(xs_test.shape)`  
`print(ys_train.shape)`  
`print(ys_test.shape)`

```
(1070,)
(268,)
(1070,)
(268,)
```

In [469]: `xs_train = xs_train.values.reshape(-1,1)`  
`xs_test = xs_test.values.reshape(-1,1)`

In [470]: `models = svm.SVC()`  
`models.fit(xs_train,ys_train)`


Out[470]: `SVC()`

In [471]: `models.score(xs_test,ys_test)`

Out[471]: 0.9328358208955224



```
In [472]: plt.scatter(x = data['charges'],y = data['smoke'])  
plt.xlabel('charges')  
plt.ylabel("smoke")  
plt.show()
```



```
In [473]: from sklearn.ensemble import RandomForestRegressor
```

```
In [476]: regressor = RandomForestRegressor(n_estimators = 500, random_state = 1)  
regressor.fit(xs_train,ys_train)
```

```
Out[476]: RandomForestRegressor(n_estimators=500, random_state=1)
```

```
In [477]: regressor.score(xs_train,ys_train)
```

```
Out[477]: 0.9414024055614973
```